

Problem statement: To predict How Best the Data Fits, To Predict the accuracy of the Rainfall based on the given features

1) Data collection ¶

```
In [1]: #Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: data=pd.read_csv(r"C:\Users\manasa\Downloads\rainfall in india 1901-2015.csv")
data
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4

4116 rows × 13 columns



2)Data Cleaning and Preprocessing

In [3]: data.head()

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4

In [4]: data.tail()

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	1
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	2

In [5]: data.shape

Out[5]: (4116, 19)

In [6]: data.describe

```
Out[6]: <bound method NDFrame.describe of
SUBDIVISION  YEAR  JAN
FEB  MAR  APR  MAY  JUN
0  ANDAMAN & NICOBAR ISLANDS  1901  49.2  87.1  29.2  2.3  528.8  517.5
\
1  ANDAMAN & NICOBAR ISLANDS  1902  0.0  159.8  12.2  0.0  446.1  537.1
2  ANDAMAN & NICOBAR ISLANDS  1903  12.7  144.0  0.0  1.0  235.1  479.9
3  ANDAMAN & NICOBAR ISLANDS  1904  9.4  14.7  0.0  202.4  304.5  495.1
4  ANDAMAN & NICOBAR ISLANDS  1905  1.3  0.0  3.3  26.9  279.5  628.7
...
4111  LAKSHADWEEP  2011  5.1  2.8  3.1  85.9  107.2  153.6
4112  LAKSHADWEEP  2012  19.2  0.1  1.6  76.8  21.2  327.0
4113  LAKSHADWEEP  2013  26.2  34.4  37.5  5.3  88.3  426.2
4114  LAKSHADWEEP  2014  53.2  16.1  4.4  14.9  57.4  244.1
4115  LAKSHADWEEP  2015  2.2  0.5  3.7  87.1  133.1  296.6

JUL  AUG  SEP  OCT  NOV  DEC  ANNUAL  Jan-Feb  Mar-May
0  365.1  481.1  332.6  388.5  558.2  33.6  3373.2  136.3  560.3  \
1  228.9  753.7  666.2  197.2  359.0  160.5  3520.7  159.8  458.3
2  728.4  326.7  339.0  181.2  284.4  225.0  2957.4  156.7  236.1
3  502.0  160.1  820.4  222.2  308.7  40.1  3079.6  24.1  506.9
4  368.7  330.5  297.0  260.7  25.4  344.7  2566.7  1.3  309.7
...
4111  350.2  254.0  255.2  117.4  184.3  14.9  1533.7  7.9  196.2
4112  231.5  381.2  179.8  145.9  12.4  8.8  1405.5  19.3  99.6
4113  296.4  154.4  180.0  72.8  78.1  26.7  1426.3  60.6  131.1
4114  116.1  466.1  132.2  169.2  59.0  62.3  1395.0  69.3  76.7
4115  257.5  146.4  160.4  165.4  231.0  159.0  1642.9  2.7  223.9

Jun-Sep  Oct-Dec
0  1696.3  980.3
1  2185.9  716.7
2  1874.0  690.6
3  1977.6  571.0
4  1624.9  630.8
...
4111  1013.0  316.6
4112  1119.5  167.1
4113  1057.0  177.6
4114  958.5  290.5
4115  860.9  555.4
```

[4116 rows x 19 columns]>

In [7]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4112 non-null   float64
3   FEB             4113 non-null   float64
4   MAR             4110 non-null   float64
5   APR             4112 non-null   float64
6   MAY             4113 non-null   float64
7   JUN             4111 non-null   float64
8   JUL             4109 non-null   float64
9   AUG             4112 non-null   float64
10  SEP             4110 non-null   float64
11  OCT             4109 non-null   float64
12  NOV             4105 non-null   float64
13  DEC             4106 non-null   float64
14  ANNUAL          4090 non-null   float64
15  Jan-Feb         4110 non-null   float64
16  Mar-May         4107 non-null   float64
17  Jun-Sep         4106 non-null   float64
18  Oct-Dec         4103 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [8]: data.isnull().sum()

```
Out[8]: SUBDIVISION      0
YEAR                  0
JAN                   4
FEB                   3
MAR                   6
APR                   4
MAY                   3
JUN                   5
JUL                   7
AUG                   4
SEP                   6
OCT                   7
NOV                  11
DEC                  10
ANNUAL               26
Jan-Feb              6
Mar-May              9
Jun-Sep             10
Oct-Dec             13
dtype: int64
```

```
In [9]: data.fillna(method="ffill",inplace=True)
```

```
In [10]: data.isnull().sum()
```


```
Out[10]: SUBDIVISION      0
YEAR                0
JAN                 0
FEB                 0
MAR                 0
APR                 0
MAY                 0
JUN                 0
JUL                 0
AUG                 0
SEP                 0
OCT                 0
NOV                 0
DEC                 0
ANNUAL              0
Jan-Feb             0
Mar-May             0
Jun-Sep             0
Oct-Dec             0
dtype: int64
```

```
In [11]: data['YEAR'].value_counts()
```


```
Out[11]: YEAR
1963      36
2002      36
1976      36
1975      36
1974      36
..
1915      35
1918      35
1954      35
1955      35
1909      34
Name: count, Length: 115, dtype: int64
```

3)Exploratory Data Analysis


```
In [12]: sns.lmplot(x='NOV',y='DEC',order=2,data=data,ci=None)  
plt.show()
```



```
In [13]: df=data[['JAN','FEB','MAR','APR','DEC']]  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```



```
In [14]: sns.pairplot(df)
plt.show()
```



4) Training our Model

```
In [15]: x=np.array(df['FEB']).reshape(-1,1)
y=x=np.array(df['JAN']).reshape(-1,1)
```


```
In [16]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30)
```

```
In [17]: lin=LinearRegression()
lin.fit(x_train,y_train)
print(lin.score(x_train,y_train))
```

1.0


5)Exploring our Results

```
In [18]: y_pred=lin.predict(x_test)
plt.scatter(x_test,y_test,color='blue')
plt.plot(x_test,y_pred,color='black')
plt.show()
```



7)Working with subset of data

```
In [19]: df700=df[:][:700]
sns.lmplot(x='FEB',y='JAN',order=2,ci=None,data=df700)
plt.show()
```



```
In [20]: df700.fillna(method='ffill',inplace=True)
```


```
In [21]: x=np.array(df700['FEB']).reshape(-1,1)
y=x=np.array(df700['JAN']).reshape(-1,1)
```

```
In [22]: df700.dropna(inplace=True)
```

```
In [24]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.03)
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

1.0

```
In [25]: y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [26]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

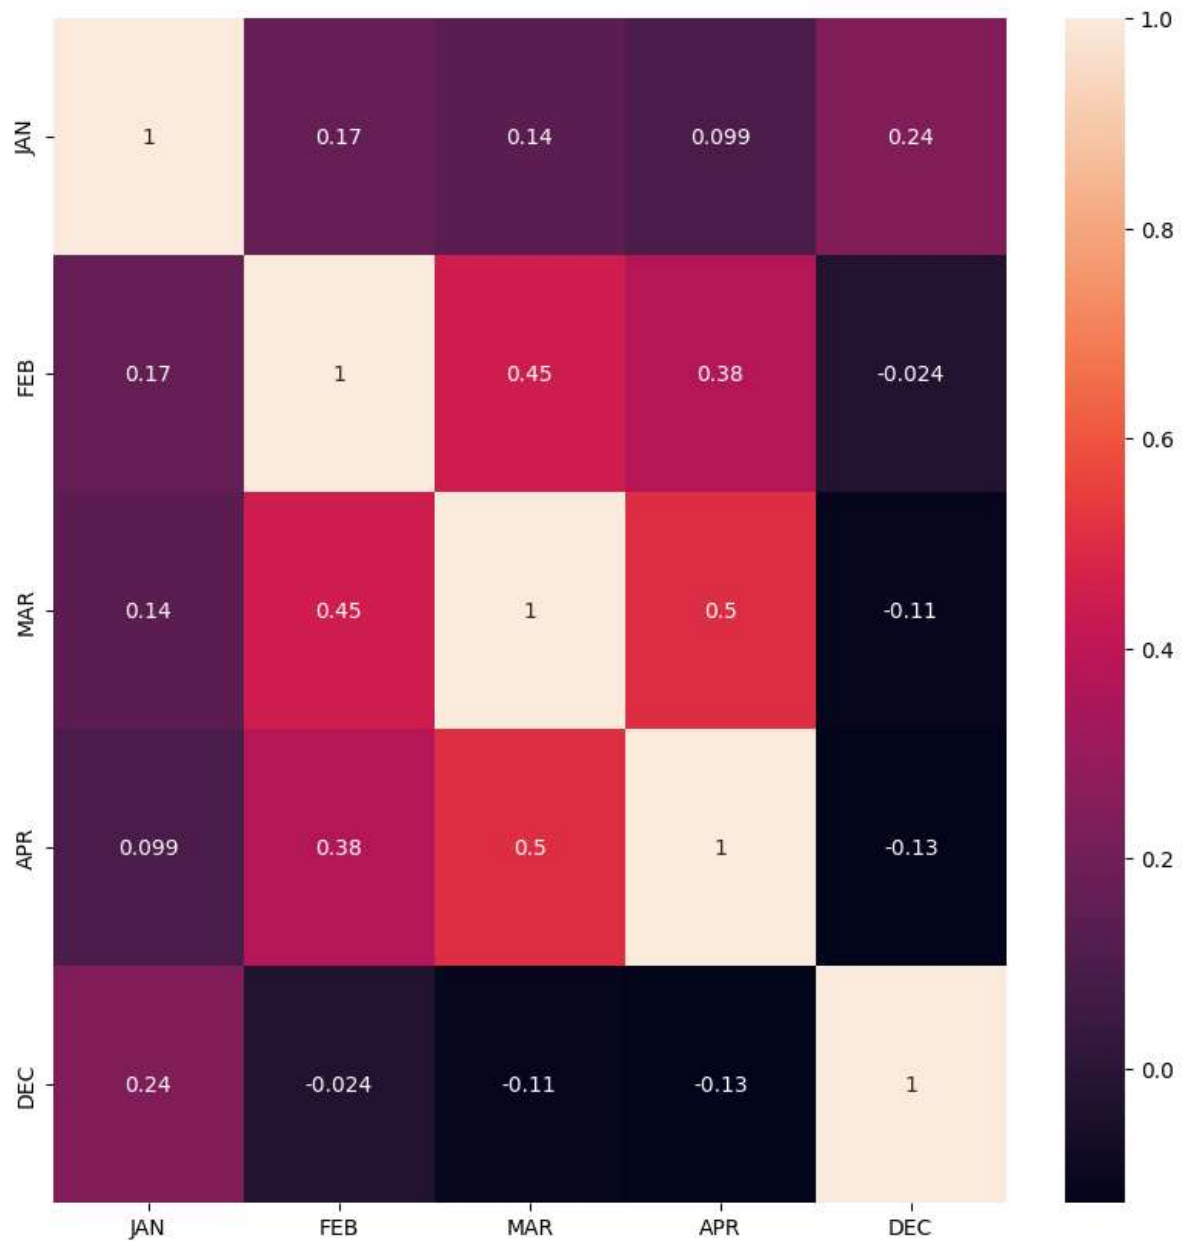
```
In [27]: lr=LinearRegression()
lr.fit(x_train,y_train)
y_pred=lr.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 1.0

Ridge Regression

```
In [28]: #Importing Libraries
from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [29]: plt.figure(figsize=(10,10))  
sns.heatmap(df700.corr(),annot=True)  
plt.show()
```



```
In [30]: features=df.columns[0:5]  
target=df.columns[-5]
```

```
In [48]: x=df[features].values  
y=df[target].values  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=42)  
print("The dimension of X_train is {}".format(x_train.shape))  
print("The dimension of X_test is {}".format(x_test.shape))
```

The dimension of X_train is (2881, 5)
The dimension of X_test is (1235, 5)

```
In [49]: lr = LinearRegression()
#Fit model
lr.fit(x_train, y_train)
#predict
actual = y_test
train_score_lr = lr.score(x_train, y_train)
test_score_lr = lr.score(x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0


```
In [51]: ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge = ridgeReg.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9999999999856335

The test score for ridge model is 0.9999999999840021

```
In [53]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=7,
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker="o",markersize=7,
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



Lasso Regression

```
In [33]: #Importing Libraries
lasso= Lasso(alpha=10)
lasso.fit(x_train,y_train)
#train and test scorefor ridge regression
train_score_ls = lasso.score(x_train, y_train)
test_score_ls= lasso.score(x_test, y_test)
print("\nLasso Model:\n")
print("The train score for lasso model is {}".format(train_score_ls))
print("The test score for lasso model is {}".format(test_score_ls))
```

Lasso Model:

The train score for lasso model is 0.9999539358821963

The test score for lasso model is 0.9999440763208948

```
In [34]: plt.figure(figsize=(10,10))
```

Out[34]: <Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

```
In [35]: from sklearn.linear_model import LassoCV
```

```
In [36]: #using the linear cv model
from sklearn.linear_model import RidgeCV
#cross validation
ridge_cv=RidgeCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
print(ridge_cv.score(x_train,y_train))
print(ridge_cv.score(x_test,y_test))
```

0.9999999999860413

0.9999999999830534

```
In [37]: #using the linear cv model
from sklearn.linear_model import LassoCV
#cross validation
lasso_cv=LassoCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.9999999999999954

0.9999999999999944

C:\Users\manasa\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model_coordinate_descent.py:1568: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

Elastic Regression

```
In [38]: from sklearn.linear_model import ElasticNet
```

```
In [39]: el=ElasticNet()  
el.fit(x_train,y_train)  
print(el.coef_)  
print(el.intercept_)  
el.score(x,y)
```

```
[0.99932152]
```

```
[0.01712925]
```

```
Out[39]: 0.9999995396414608
```

```
In [40]: y_pred_elastic=el.predict(x_train)
```

```
In [41]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)  
print(mean_squared_error)
```

```
2944.787483507373
```

CONCLUSION:

The given data is "Rain fall prediction".here we need to find the best fit model. As per the given data set I had applied different types of models...in which different type of models got different type of accuracies. The accuracy of the Linear Regression is 1.0. The accuracy of the Ridge Model is 0.9999999999856. The accuracy of the Lasso Model is 0.20. The accuracy of the ElasticNet Regression is 0.99999914, comparing to all the models, Ridge Regression got the Highest Accuracy

```
In [ ]:
```